

Question 8.1

Using the same crime data set `uscrime.txt` as in Question 8.2, apply Principal Component Analysis and then create a regression model using the first few principal components. Specify your new model in terms of the original variables (not the principal components), and compare its quality to that of your solution to Question 8.2. You can use the R function `prcomp` for PCA. (Note that to first scale the data, you can include `scale. = TRUE` to scale as part of the PCA function. Don't forget that, to make a prediction for the new city, you'll need to unscale the coefficients (i.e., do the scaling calculation in reverse)!)

Approach Followed

1) After data is ingested, I followed series of steps to see the co-correlation of the predictors among themselves and its impact on the Crime responses using Correlation plots, Correlation matrix and the Correlation tables. In my opinion Correlation tables were the best approach to quickly see the correlated fields. 2) This step is followed by creating eigen values and eigen vectors for $X(\text{Transpose})^*X$; X being the matrix of the Crime data set. I used Caret transformation to get predictors, box cox transformation and find principal components which suggested that "PCA needed 9 components to capture 95 percent of the variance" 3) Based on this recommendation I set my "n" values for PCA as 9. Using `prcomp`, I created PCA and obtained first "9" components using which I created linear regression model (using `lm`). The model gave the adjusted R squared :0.69; R squared :0.61

4) Now, per the requirement, we have to express the model in terms of original variables and NOT principal components. To achieve this I obtained my components PC1 through PC9 (and then used `pca rotation and transform` function to get the original Coefficients. One item to notice is since the coefficients are scaled earlier, I need to unscale it to get the exact coefficients of the model. 5) To scale, it subtracts the mean and divides by the standard deviation, for each variable. I used this logic to unscale the coefficients 6) Compare the PCA model accuracy vs. the original model's accuracy. The Original model R2 squared is 0.80 and Adjusted R2 .70 which is more than the PCA model accuracy. Then I ran a loop from 1 through all PCA components to see any of its accuracy beat the original model.. But none. As noted in the lectures, overfitting could have occurred with less data points with original linear regression model 7) Finally the prediction for the new City using `predict` function. I used both the model to predict the outcome of Crime. Predicted value of new data using PCR is 1112 vs. 155 crime data value using original regression with all the predictors

#PCA generated crime ratio 7 times more than the linear regression model with all predictors for the new data value

```
In [203]: # Clear environment
rm(list = ls())
# Setting the random number generator seed so that our results are reproducible
set.seed(1)
```

```
In [204]: #####LIBRARY#####
library("ggplot2")
#install.packages("devtools")
#install.packages("corrplot")
library("devtools")
#library("corrplot") # to use Correlation plot
library(tidyr) # to use "as tibble,"
library(plyr) #to use arrange function
library(car) #Scatterplotmatrix
library(GGally) #for ggpairs graph
library(pls)# to get pca fit
```

INGESTION OF DATA AND STUDY THE RELATIONSHIPS

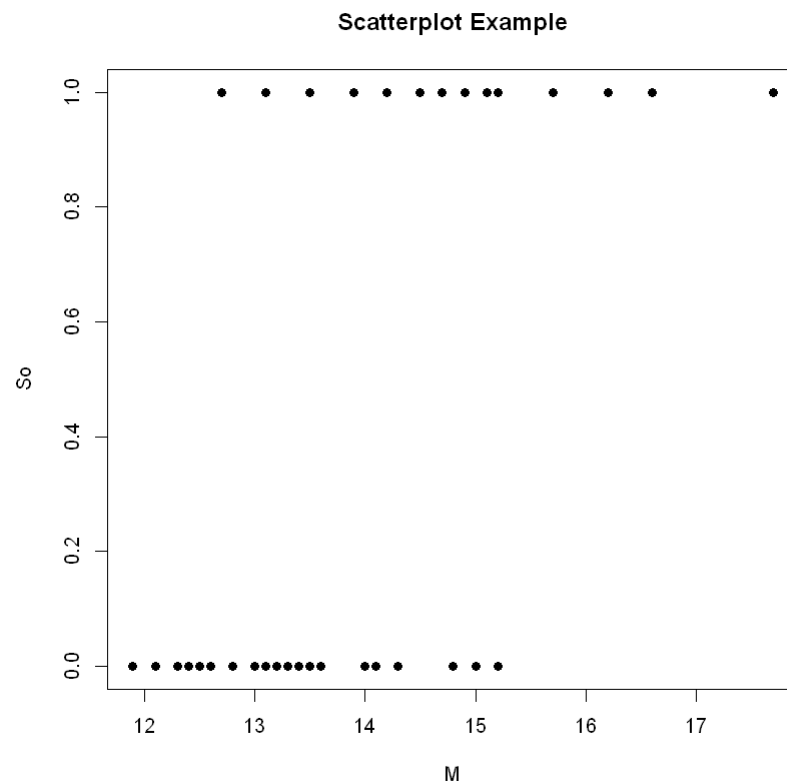
```
In [205]: #####INGEST FILE#####


#crime<- read.table("C:\\Preethi\\R\\USCrimes.txt",header=TRUE,stringsAsFactors = FALSE,sep="\t")
data = read.table("C:\\Preethi\\R\\USCrimes.txt",header=TRUE,stringsAsFactors = FALSE,sep="\t") %>% as.data.frame()
head(data,2)
```


A data.frame: 2 × 16

	M	So	Ed	Po1	Po2	LF	M.F	Pop	NW	U1	U2	Wealth	Ineq	Prob	Time	Crime
	<dbl>	<int>	<dbl>	<dbl>	<dbl>	<dbl>	<dbl>	<int>	<dbl>	<dbl>	<dbl>	<int>	<dbl>	<dbl>	<dbl>	<int>
1	15.1	1	9.1	5.8	5.6	0.510	95.0	33	30.1	0.108	4.1	3940	26.1	0.084602	26.2011	791
2	14.3	0	11.3	10.3	9.5	0.583	101.2	13	10.2	0.096	3.6	5570	19.4	0.029599	25.2999	1635

```
In [225]: ##EXPLORING DATA#####  
#2-D graphs for each predictors to understand relationship and correlation of predictors  
#Just reduced to 2 charts to avoid overfill of assignment with scatterplot  
for (i in 1:3){  
  for (j in 1:3){  
    if (i<j){  
      plot(data[,i],data[,j], main="Scatterplot Example",xlab=colnames(data)[i],ylab=colnames(data)[j], pch=19)  
    }  
  }  
}
```

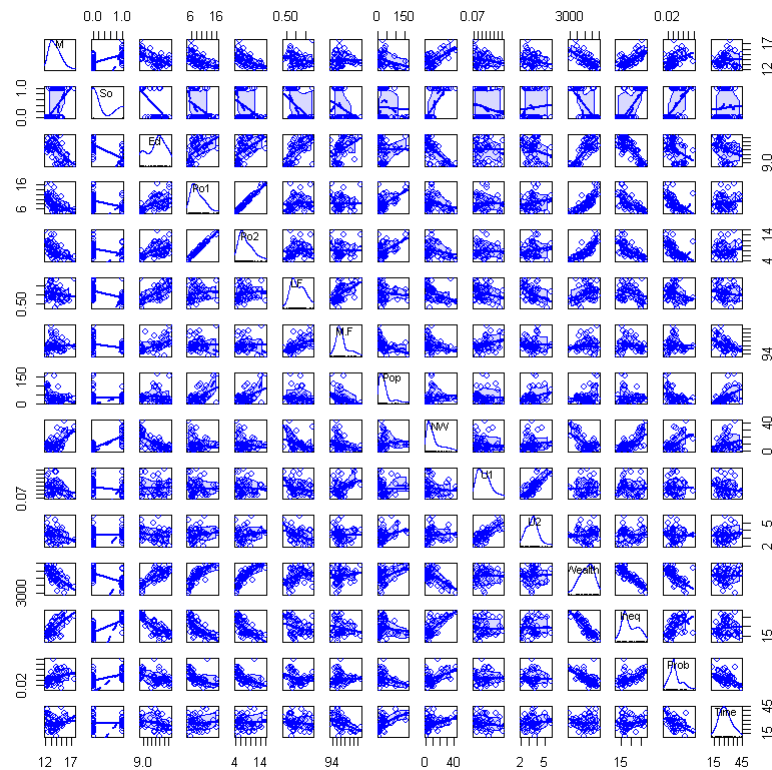




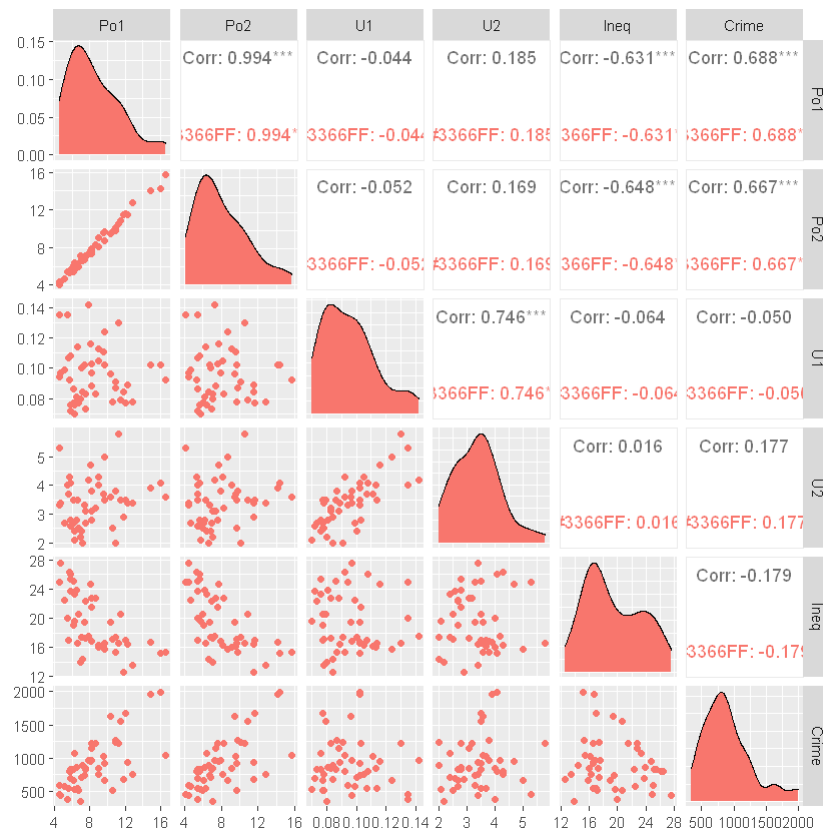


```
In [144]: #Library(car)
#ALL predictors scatterplot
scatterplotMatrix(~ M+So+Ed+Po1+Po2+LF+M.F+Pop+NW+U1+U2+Wealth+Ineq+Prob+Time, data = data)
#too tough to read
```

Warning message in smoother(x[subs], y[subs], col = smoother.args\$col[i], log.x = FALSE, :
"could not fit smooth"



```
In [141]: #USING SUBSET OF PREDICTORS
ggpairs(data, columns = c("Po1", "Po2", "U1", "U2", "Ineq", "Crime"),
        mapping=ggplot2::aes(color= "#3366FF"))
```



```
In [142]: #STUDY THE CORRELATION OF THE PREDICTORS (CLOSER TO 1, MORE CORRELATED THE PREDICTORS ARE..)
corr <- cor(data)
round(corr, 2)
#Observation co-relation between predictors
# high POSITIVE Correlation
##Po1-Po2;U1-U2;Po1/Po2-Wealth;So-NW;So-Ineq;Ed-Wealth
# high NEGATIVE Correlation
## So-Ed;Ed-Ineq;P
```

A matrix: 16 × 16 of type dbl

	M	So	Ed	Po1	Po2	LF	M.F	Pop	NW	U1	U2	Wealth	Ineq	Prob	Time	Crime
M	1.00	0.58	-0.53	-0.51	-0.51	-0.16	-0.03	-0.28	0.59	-0.22	-0.24	-0.67	0.64	0.36	0.11	-0.09
So	0.58	1.00	-0.70	-0.37	-0.38	-0.51	-0.31	-0.05	0.77	-0.17	0.07	-0.64	0.74	0.53	0.07	-0.09
Ed	-0.53	-0.70	1.00	0.48	0.50	0.56	0.44	-0.02	-0.66	0.02	-0.22	0.74	-0.77	-0.39	-0.25	0.32
Po1	-0.51	-0.37	0.48	1.00	0.99	0.12	0.03	0.53	-0.21	-0.04	0.19	0.79	-0.63	-0.47	0.10	0.69
Po2	-0.51	-0.38	0.50	0.99	1.00	0.11	0.02	0.51	-0.22	-0.05	0.17	0.79	-0.65	-0.47	0.08	0.67
LF	-0.16	-0.51	0.56	0.12	0.11	1.00	0.51	-0.12	-0.34	-0.23	-0.42	0.29	-0.27	-0.25	-0.12	0.19
M.F	-0.03	-0.31	0.44	0.03	0.02	0.51	1.00	-0.41	-0.33	0.35	-0.02	0.18	-0.17	-0.05	-0.43	0.21
Pop	-0.28	-0.05	-0.02	0.53	0.51	-0.12	-0.41	1.00	0.10	-0.04	0.27	0.31	-0.13	-0.35	0.46	0.34
NW	0.59	0.77	-0.66	-0.21	-0.22	-0.34	-0.33	0.10	1.00	-0.16	0.08	-0.59	0.68	0.43	0.23	0.03
U1	-0.22	-0.17	0.02	-0.04	-0.05	-0.23	0.35	-0.04	-0.16	1.00	0.75	0.04	-0.06	-0.01	-0.17	-0.05
U2	-0.24	0.07	-0.22	0.19	0.17	-0.42	-0.02	0.27	0.08	0.75	1.00	0.09	0.02	-0.06	0.10	0.18
Wealth	-0.67	-0.64	0.74	0.79	0.79	0.29	0.18	0.31	-0.59	0.04	0.09	1.00	-0.88	-0.56	0.00	0.44
Ineq	0.64	0.74	-0.77	-0.63	-0.65	-0.27	-0.17	-0.13	0.68	-0.06	0.02	-0.88	1.00	0.47	0.10	-0.18
Prob	0.36	0.53	-0.39	-0.47	-0.47	-0.25	-0.05	-0.35	0.43	-0.01	-0.06	-0.56	0.47	1.00	-0.44	-0.43
Time	0.11	0.07	-0.25	0.10	0.08	-0.12	-0.43	0.46	0.23	-0.17	0.10	0.00	0.10	-0.44	1.00	0.15
Crime	-0.09	-0.09	0.32	0.69	0.67	0.19	0.21	0.34	0.03	-0.05	0.18	0.44	-0.18	-0.43	0.15	1.00

CHECKING THEORY : EIGEN VALUES AND VECTORS

```
In [23]: #Eigen values & vectors
X <- as.matrix(data)
X_Trans_X <- t(X)%*%X
ev <- eigen(X_Trans_X)
#Eigenvalues
ev$values
#Eigenvectors
#ev$vectors
```

```
1380740012.90473 · 5384801.42652306 · 59330.4704502898 · 20626.7567816144 · 2187.41595483789 ·
1677.27232653165 · 151.07239653922 · 88.021838531293 · 41.095517797618 · 24.8002928942946 ·
9.65667653420426 · 2.80155722968799 · 1.96793750526589 · 0.0256320722571422 · 0.00769541551602604 ·
0.00252687018876677
```

```
In [136]: #Find Determinant of X(Transpose)*X -lambda(identity matrix) for 1st Eigen value and vector and test the theory
#diag(ncol(data))
identity_mat <-diag(ncol(data))
det(X_Trans_X-ev$values[2]*identity_mat)
```

```
4.25568663425159e+94
```

```
In [206]: #Caret to transform predictors, box cox tranformation and find princial components
pca_data <-caret::preProcess(
  data %>% dplyr::select(-Crime),
  method=c('center','scale','nzv','pca'))
pca_data # suggestion is to use PCA for 9 components out of 15 components
#linear combination of predictors for each principal component
pca_data$rotation
```

Created from 47 samples and 15 variables

Pre-processing:

- centered (15)
- ignored (0)
- principal component signal extraction (15)
- scaled (15)

PCA needed 9 components to capture 95 percent of the variance

A matrix: 15 × 9 of type dbl

	PC1	PC2	PC3	PC4	PC5	PC6	PC7	PC8	PC9
M	-0.30371194	0.06280357	0.1724199946	-0.02035537	-0.35832737	-0.449132706	-0.15707378	-0.55367691	0.154747
So	-0.33088129	-0.15837219	0.0155433104	0.29247181	-0.12061130	-0.100500743	0.19649727	0.22734157	-0.65599
Ed	0.33962148	0.21461152	0.0677396249	0.07974375	-0.02442839	-0.008571367	-0.23943629	-0.14644678	-0.44326
Po1	0.30863412	-0.26981761	0.0506458161	0.33325059	-0.23527680	-0.095776709	0.08011735	0.04613156	0.194254
Po2	0.31099285	-0.26396300	0.0530651173	0.35192809	-0.20473383	-0.119524780	0.09518288	0.03168720	0.195120
LF	0.17617757	0.31943042	0.2715301768	-0.14326529	-0.39407588	0.504234275	-0.15931612	0.25513777	0.143934
M.F	0.11638221	0.39434428	-0.2031621598	0.01048029	-0.57877443	-0.074501901	0.15548197	-0.05507254	-0.24378
Pop	0.11307836	-0.46723456	0.0770210971	-0.03210513	-0.08317034	0.547098563	0.09046187	-0.59078221	-0.20244
NW	-0.29358647	-0.22801119	0.0788156621	0.23925971	-0.36079387	0.051219538	-0.31154195	0.20432828	0.189841
U1	0.04050137	0.00807439	-0.6590290980	-0.18279096	-0.13136873	0.017385981	-0.17354115	-0.20206312	0.020693
U2	0.01812228	-0.27971336	-0.5785006293	-0.06889312	-0.13499487	0.048155286	-0.07526787	0.24369650	0.055760
Wealth	0.37970331	-0.07718862	0.0100647664	0.11781752	0.01167683	-0.154683104	-0.14859424	0.08630649	-0.23196

	PC1	PC2	PC3	PC4	PC5	PC6	PC7	PC8	PC9
Ineq	-0.36579778	-0.02752240	-0.0002944563	-0.08066612	-0.21672823	0.272027031	0.37483032	0.07184018	-0.02494
Prob	-0.25888661	0.15831708	-0.1176726436	0.49303389	0.16562829	0.283535996	-0.56159383	-0.08598908	-0.05306
Time	-0.02062867	-0.38014836	0.2235664632	-0.54059002	-0.14764767	-0.148203050	-0.44199877	0.19507812	-0.23551

PCA COMPUTATION USING PRCOMP

```
In [207]: pca <- prcomp(data[,1:15], scale. = TRUE)
          summary(pca)
```

Importance of components:

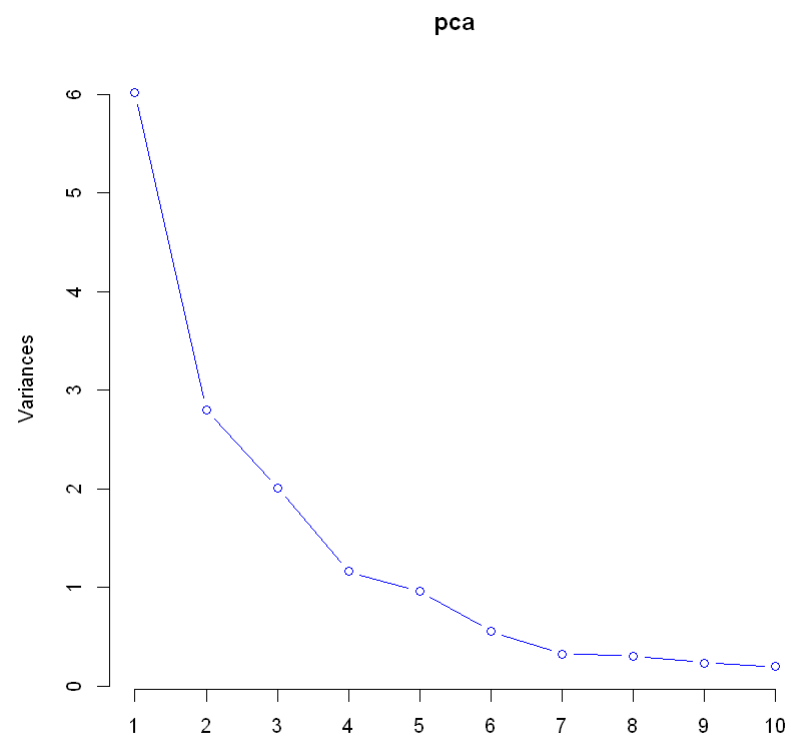
```

              PC1    PC2    PC3    PC4    PC5    PC6    PC7
Standard deviation  2.4534 1.6739 1.4160 1.07806 0.97893 0.74377 0.56729
Proportion of Variance 0.4013 0.1868 0.1337 0.07748 0.06389 0.03688 0.02145
Cumulative Proportion 0.4013 0.5880 0.7217 0.79920 0.86308 0.89996 0.92142

              PC8    PC9    PC10    PC11    PC12    PC13    PC14
Standard deviation  0.55444 0.48493 0.44708 0.41915 0.35804 0.26333 0.2418
Proportion of Variance 0.02049 0.01568 0.01333 0.01171 0.00855 0.00462 0.0039
Cumulative Proportion 0.94191 0.95759 0.97091 0.98263 0.99117 0.99579 0.9997

              PC15
Standard deviation  0.06793
Proportion of Variance 0.00031
Cumulative Proportion 1.00000
```

```
In [150]: plot(pca, type="lines",col="blue")
```



```
In [208]: #get first 9 Principal componet
PCs <- pca$x[,1:9]
head(PCs,5)
```

A matrix: 5 × 9 of type dbl

PC1	PC2	PC3	PC4	PC5	PC6	PC7	PC8	PC9
-4.199284	-1.0938312	-1.11907395	0.67178115	0.05528338	0.3073383	-0.56640816	-0.007801727	0.22350995
1.172663	0.6770136	-0.05244634	-0.08350709	-1.17319982	-0.5832373	0.19561119	0.154566472	0.43677720
-4.173725	0.2767750	-0.37107658	0.37793995	0.54134525	0.7187223	0.10330693	0.351138883	0.06299232
3.834962	-2.5769060	0.22793998	0.38262331	-1.64474650	0.7294884	0.26699499	-1.547460841	-0.37954181
1.839300	1.3309856	1.27882805	0.71814305	0.04159032	-0.3940902	0.07050766	-0.543237437	0.22463245

Build linear regression model with the first 9 principal components

```
In [209]: PCcrime <- cbind(PCs, data$Crime) #Create new data matrix with first 5 PCs and crime rate
#PCcrime
model <- lm(V10~., data = as.data.frame(PCcrime)) #Create regression model based on 5 principal components
#model
summary(model) #PCA: Adjusted R squared :0.69;R squared :0.61
```

Call:

```
lm(formula = V10 ~ ., data = as.data.frame(PCcrime))
```

Residuals:

Min	1Q	Median	3Q	Max
-455.9	-132.5	21.5	139.9	393.0

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)	
(Intercept)	905.09	34.91	25.928	< 2e-16	***
PC1	65.22	14.38	4.535	5.88e-05	***
PC2	-70.08	21.08	-3.325	0.00201	**
PC3	25.19	24.92	1.011	0.31857	
PC4	69.45	32.73	2.122	0.04061	*
PC5	-229.04	36.04	-6.355	2.08e-07	***
PC6	-60.21	47.44	-1.269	0.21228	
PC7	117.26	62.20	1.885	0.06728	.
PC8	28.72	63.64	0.451	0.65446	
PC9	-37.18	72.76	-0.511	0.61244	

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 239.3 on 37 degrees of freedom

Multiple R-squared: 0.692, Adjusted R-squared: 0.6171

F-statistic: 9.239 on 9 and 37 DF, p-value: 3.588e-07

Specify the new model in terms of the original variables (not the principal components), UNSCALE data in revrse

```
In [212]: # Get coefficients in terms of original data from PCA coefficients
model$coefficient
beta0 <- model$coefficients[1] #intercept
betas <- model$coefficients[2:10] #PC
print(paste("Intercept Co-efficient: ",beta0))
betas
```

(Intercept): 905.085106382979 **PC1:** 65.215930138666 **PC2:** -70.0831185497858 **PC3:** 25.1940780425772 **PC4:** 69.446030796839
PC5: -229.042822001686 **PC6:** -60.2132861756709 **PC7:** 117.255897957602 **PC8:** 28.7165560702805 **PC9:** -37.1756419454256

[1] "Intercept Co-efficient: 905.085106382979"

PC1: 65.215930138666 **PC2:** -70.0831185497858 **PC3:** 25.1940780425772 **PC4:** 69.446030796839 **PC5:** -229.042822001686 **PC6:** -60.2132861756709 **PC7:** 117.255897957602 **PC8:** 28.7165560702805 **PC9:** -37.1756419454256

```
In [213]: # Transform the PC coefficients into coefficients for the original variables
```

```
#pca$rotation[,1:4]
alphas <- pca$rotation[,1:9] %*% betas
t(alphas) #transform PC co-efficients into original variable co-efficients
```

A matrix: 1 × 15 of type dbl

M	So	Ed	Po1	Po2	LF	M.F	Pop	NW	U1	U2	Wealth	Ineq
47.76773	97.85583	4.661941	126.6093	123.4647	29.18835	138.325	27.10574	57.26768	-25.10066	25.54553	38.92574	52.665



```
In [214]: # Above is SCALED coefficients
#To convert into original
# When scaling, this function subtracts the mean and divides by the standard deviation, for each variable.
# So,  $\alpha * (x - \text{mean}) / \text{sd} = \text{originalAlpha} * x$ .
# That means:
# (1)  $\text{originalAlpha} = \alpha / \text{sd}$ 
# (2) we have to modify the constant term  $a_0$  by  $\alpha * \text{mean} / \text{sd}$ 

originalAlpha <- alphas/sapply(data[,1:15],sd)
originalBeta0 <- beta0 - sum(alphas*sapply(data[,1:15],mean)/sapply(data[,1:15],sd))
originalBeta0
```

(Intercept): -5742.13576339317

```
In [216]: # Here are the coefficients for unscaled data:
t(originalAlpha)
```

A matrix: 1 × 15 of type dbl

M	So	Ed	Po1	Po2	LF	M.F	Pop	NW	U1	U2	Wealth	In
38.00853	204.3025	4.167285	42.60219	44.15555	722.2726	46.94177	0.7119751	5.569224	-1392.255	30.24768	0.04034134	13

In [217]: *#ORIGINAL MODEL Now let's compare with the regression model from the previous homework*

```
model2 <- lm( Crime ~ ., data = data)
summary(model2) #R2 .80 and Adju R2 .70 (vs.#PCA: Adjusted R squared :0.69;R squared :0.61)
```

Call:

```
lm(formula = Crime ~ ., data = data)
```

Residuals:

	Min	1Q	Median	3Q	Max
	-395.74	-98.09	-6.69	112.99	512.67

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)	
(Intercept)	-5.984e+03	1.628e+03	-3.675	0.000893	***
M	8.783e+01	4.171e+01	2.106	0.043443	*
So	-3.803e+00	1.488e+02	-0.026	0.979765	
Ed	1.883e+02	6.209e+01	3.033	0.004861	**
Po1	1.928e+02	1.061e+02	1.817	0.078892	.
Po2	-1.094e+02	1.175e+02	-0.931	0.358830	
LF	-6.638e+02	1.470e+03	-0.452	0.654654	
M.F	1.741e+01	2.035e+01	0.855	0.398995	
Pop	-7.330e-01	1.290e+00	-0.568	0.573845	
NW	4.204e+00	6.481e+00	0.649	0.521279	
U1	-5.827e+03	4.210e+03	-1.384	0.176238	
U2	1.678e+02	8.234e+01	2.038	0.050161	.
Wealth	9.617e-02	1.037e-01	0.928	0.360754	
Ineq	7.067e+01	2.272e+01	3.111	0.003983	**
Prob	-4.855e+03	2.272e+03	-2.137	0.040627	*
Time	-3.479e+00	7.165e+00	-0.486	0.630708	

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 209.1 on 31 degrees of freedom

Multiple R-squared: 0.8031, Adjusted R-squared: 0.7078

F-statistic: 8.429 on 15 and 31 DF, p-value: 3.539e-07

Compare quality of PCA Model vs. Original Model based on straightforward regression model

```
In [ ]: # These results suggest that we are better off using a more straightforward regression model
# instead of PCA before using regression.
# If we had used all 15 principal components, we would have obtained
# an R-squared value of 0.803, which is the same R-squared value when using all
# 15 regular predictors in a basic linear regression model.
```

```
In [166]: # all possibilities: for i=1..15, run a regression using the first i principal components
r2 <- numeric(15) # create a vector to store the R-squared values
for (i in 1:15) {
  pclist <- pca$x[,1:i] # use the first i principal components
  pcc <- cbind(data[,16],pclist) # create data set
  model <- lm(V1~.,data = as.data.frame(pcc)) # fit model
  r2[i] <- 1 - sum(model$residuals^2)/sum((data$Crime - mean(data$Crime))^2) # calculate R-squared
}
r2
```

```
0.171135123438665 · 0.263133909632497 · 0.271641569122873 · 0.309112068654877 · 0.645194053656692 ·
0.658602327815706 · 0.688181873005123 · 0.689876527649446 · 0.692049141529042 · 0.696287251226172 ·
0.697386539400947 · 0.769265609989774 · 0.77236636338524 · 0.791144655274498 · 0.803086758316909
```

```
In [167]: # All PCA first "n" components generate model R-square less than original model R square value of 0.80.
#As noted earlier, with smaller subset of data, we are running into overfitting yielding high R square for origin
```

Predict new observations' Crime fit

crime prediction for new point of observation based on Original vs PCA model

```
In [218]: #datapoint from previous homework
new_observe = data.frame(
  M = 14.0,
  So = 0,
  Ed = 10.0,
  Po1 = 12.0,
  Po2 = 15.5,
  LF = 0.640,
  M.F = 94.0 ,
  Pop = 150,
  NW = 1.1,
  U1 = 0.120,
  U2 = 3.6 ,
  Wealth = 3200,
  Ineq = 20.1 ,
  Prob = 0.04 ,
  Time = 39.0
)
```

```
In [219]: #ORIGINAL Prediction
#model2 <- lm( Crime ~ ., data = data)
crime_pred_orig = predict(model2, new_observe) %>%
  as_tibble()
crime_pred_orig
```

A tibble: 1

× 1

value
<dbl>
155.4349

```
In [221]: #PCA FITTED PREDICTION
library(pls)# to get pcr fit
# Run principal component regression function with only the first 4 principal components

numcomp <- 9
pcr.fit <- pcr(Crime ~ ., data = data, scale = TRUE, ncomp = numcomp)
summary(pcr.fit)
head(pcr.fit$scores,5)
```

Data: X dimension: 47 15

Y dimension: 47 1

Fit method: svdpc

Number of components considered: 9

TRAINING: % variance explained

	1 comps	2 comps	3 comps	4 comps	5 comps	6 comps	7 comps	8 comps
X	40.13	58.81	72.17	79.92	86.31	90.00	92.14	94.19
Crime	17.11	26.31	27.16	30.91	64.52	65.86	68.82	68.99
	9 comps							
X	95.76							
Crime	69.20							

A matrix: 5 × 9 of type dbl

	Comp 1	Comp 2	Comp 3	Comp 4	Comp 5	Comp 6	Comp 7	Comp 8	Comp 9
1	-4.199284	-1.0938312	-1.11907395	0.67178115	0.05528338	0.3073383	-0.56640816	-0.007801727	0.22350995
2	1.172663	0.6770136	-0.05244634	-0.08350709	-1.17319982	-0.5832373	0.19561119	0.154566472	0.43677720
3	-4.173725	0.2767750	-0.37107658	0.37793995	0.54134525	0.7187223	0.10330693	0.351138883	0.06299232
4	3.834962	-2.5769060	0.22793998	0.38262331	-1.64474650	0.7294884	0.26699499	-1.547460841	-0.37954181
5	1.839300	1.3309856	1.27882805	0.71814305	0.04159032	-0.3940902	0.07050766	-0.543237437	0.22463245

```
In [222]: #use model to make predictions on a test set
test <- pcr(Crime ~ ., data = data, scale=TRUE, validation="CV") #Cross validation
summary(test)
#Root Mean Square error(RMSE)
#If we use interceptonly , the square error is 390.9
#If we use first components , the square error is dropped to 364
#If we add in second components , the square error is dropped to 354
```

Data: X dimension: 47 15
 Y dimension: 47 1
 Fit method: svdpc
 Number of components considered: 15

VALIDATION: RMSEP

Cross-validated using 10 random segments.

	(Intercept)	1 comps	2 comps	3 comps	4 comps	5 comps	6 comps
CV	390.9	364.8	354.1	364.0	366.9	262.5	262.8
adjCV	390.9	364.0	352.6	362.2	367.0	260.3	260.4

	7 comps	8 comps	9 comps	10 comps	11 comps	12 comps	13 comps
CV	259.6	266.7	270.6	276.6	298.8	264.5	276.4
adjCV	253.7	263.7	267.7	273.2	298.6	260.1	271.9

	14 comps	15 comps
CV	271.2	266.2
adjCV	265.7	261.0

TRAINING: % variance explained

	1 comps	2 comps	3 comps	4 comps	5 comps	6 comps	7 comps	8 comps
X	40.13	58.81	72.17	79.92	86.31	90.00	92.14	94.19
Crime	17.11	26.31	27.16	30.91	64.52	65.86	68.82	68.99

	9 comps	10 comps	11 comps	12 comps	13 comps	14 comps	15 comps
X	95.76	97.09	98.26	99.12	99.58	99.97	100.00
Crime	69.20	69.63	69.74	76.93	77.24	79.11	80.31

```
In [223]: pcr_pred <- predict(test, new_observe, ncomp=4)
pcr_pred
# predicted value of new data using PCR is 1112 vs. 155 crime data value using original regression with all the
#PCA generated crime ratio 7 times more than the linear regression model with all predictors for the new data val
```

1112.67763659883

In []: #####THE END of PCA REGRESSION ASSIGNMENT#####